

Introduction to Residual Neural Network (ResNet)

Lu Lu

Jan 25, 2019 @Crunch Seminar

Overview

- ▶ Background
- ▶ Residual neural network
- ▶ Variants of residual blocks
- ▶ Some analysis

Why Deep Neural Networks?

- ▶ Shallow NNs (single hidden layer)
 - ▶ ¹Universal approximation theorem (uniformly)
 - ▶ ² $\epsilon^{-d/n}$ neurons can approximate any C^n -function on a compact set in \mathbb{R}^d with error ϵ
- ▶ Deep NNs: better than shallow NNs (of comparable size)
 - ▶ ³Exists a function expressible by a 3-layer NN, which cannot be approximated by any 2-layer network
 - ▶ ⁴ $\frac{\text{size}_{\text{deep}}}{\text{size}_{\text{shallow}}} \approx \epsilon^d$

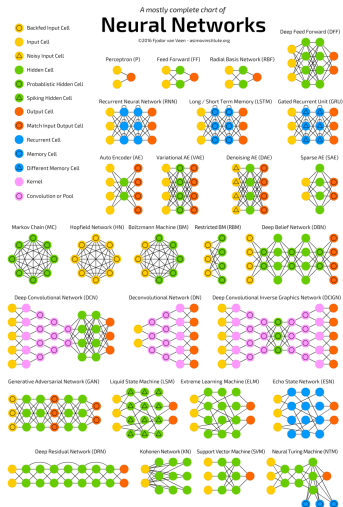
¹Cybenko, Math. Control Signals Syst., 1989; Hornik et al., Neural Netw., 1989.

²Mhaskar, Neural Comput., 1996.

³Eldan & Shamir, COLT, 2016.

⁴Mhaskar & Poggio, Anal. Appl., 2016; Mhaskar et al., AAI, 2017; Poggio et al., IJAC, 2017.

“In theory, theory and practice are the same. In practice, they are not.” — Albert Einstein?



Overview

- ▶ Background
- ▶ Residual neural network
- ▶ Variants of residual blocks
- ▶ Some analysis

Residual Neural Network⁵

- ▶ Deep networks are hard to train: vanishing gradients
- ▶ Core idea: “identity shortcut connection” that skips one or more layers
- ▶ Widely used: simple & powerful


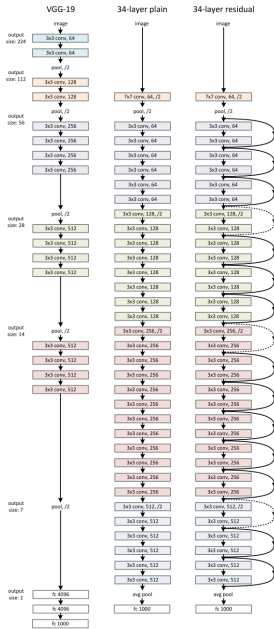


Figure : A residual block

⁵He et al., CVPR, 2016.

Residual Neural Network



Residual Neural Network

- ▶ $\mathcal{H}(\mathbf{x}) = \mathcal{F}(\mathbf{x}) + \mathbf{x} \Rightarrow \mathcal{F}(\mathbf{x}) = \mathcal{H}(\mathbf{x}) - \mathbf{x}$ (residual)
- ▶ *Hypotheses*: the residual may be an easier function to fit

\mathcal{F} ?

- ▶ If \mathcal{F} has two layers, $\mathcal{F}(\mathbf{x}) = W_2\sigma(W_1\mathbf{x})$
- ▶ If \mathcal{F} has one layers, $\mathcal{F}(\mathbf{x}) = W_1\mathbf{x}$,
 $\mathcal{H}(\mathbf{x}) = W_1\mathbf{x} + \mathbf{x} = (W_1 + 1)\mathbf{x}$. No advantage!


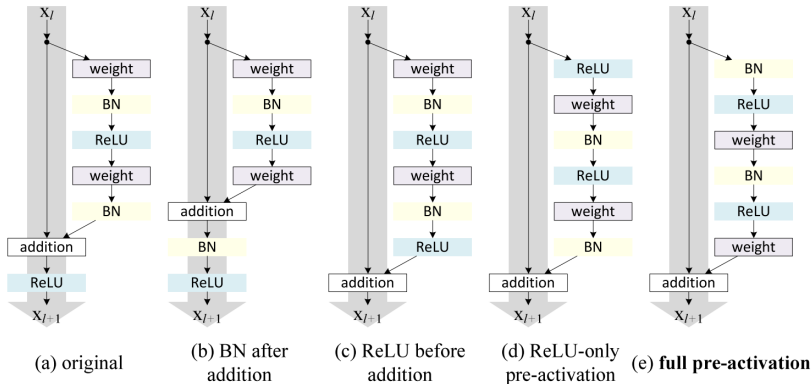


Figure : A residual block

Overview

- ▶ Background
- ▶ Residual neural network
- ▶ Variants of residual blocks
- ▶ Some analysis

Variants of Residual Blocks⁶



⁶He et al., ECCV, 2016.

Variants of Residual Blocks

Accuracy can be gained more efficiently by increasing the cardinality than by going deeper or wider.

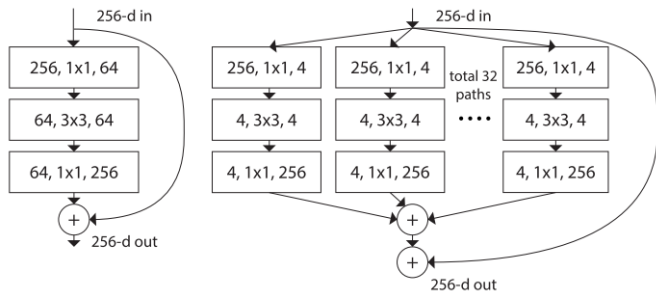


Figure : ResNeXt⁷: split-transform-merge

⁷Xie et al., CVPR, 2017.

Variants of Residual Blocks




Figure : DenseNet⁸

⁸Huang et al., CVPR, 2017.


Overview

- ▶ Background
- ▶ Residual neural network
- ▶ Variants of residual blocks
- ▶ Some analysis

Unraveled View⁹


$$\begin{aligned}y_3 &= y_2 + f_3(y_2) \\&= [y_1 + f_2(y_1)] + f_3(y_1 + f_2(y_1)) \\&= [y_0 + f_1(y_0) + f_2(y_0 + f_1(y_0))] + f_3(y_0 + f_1(y_0) + f_2(y_0 + f_1(y_0)))\end{aligned}$$

2^n paths connecting input to output layers



(a) Conventional 3-block residual network

=




(b) Unraveled view of (a)

⁹Veit et al., NIPS, 2016.

Ensemble-like Behavior¹⁰


- ▶ Lesion study: randomly deleting several modules
- ▶ Paths do not strongly depend on each other



¹⁰Veit et al., NIPS, 2016.

Vanishing Gradients?¹¹

- ▶ The effective paths are relatively shallow
- ▶ Only the short paths contribute gradients
- ▶ ResNet does not resolve vanishing gradients by preserving gradient flow throughout the entire network. Rather, they enable very deep networks by *shortening the effective paths*.



¹¹Veit et al., NIPS, 2016.

Universal Approximation

Recall¹²:

- ▶ To approximate any continuous function $[0, 1]^d \rightarrow \mathbb{R}$ by ReLU NN: minimal width is $d + 1$

ResNet with one hidden neuron:

$$\mathcal{H}(\mathbf{x}) = \mathbf{W}_2 \text{ReLU}(\mathbf{W}_1 \cdot \mathbf{x} + b) + \mathbf{x}$$

- ▶ Identity map (d dim) + one hidden neuron = $d + 1$ units


ResNet with one neuron per hidden layer: universal approximation (in L^1) for any Lebesgue-integrable function as the depth $\rightarrow \infty$.¹³

¹²Hanin et al., arXiv, 2017.

¹³Lin & Jegelka, NIPS, 2018.

Why Easier to Train?

- ▶ 2D visualization of the loss surface by “filter normalization” method¹⁴
- ▶ BoostResNet¹⁵: a training algorithm (non-differentiable), training error decays exponentially with depth



¹⁴Li et al., NIPS, 2018.

¹⁵Huang et al., ICML, 2018.

Overview

- ▶ Background
- ▶ Residual neural network
- ▶ Variants of residual blocks
- ▶ Some analysis

Cerebral Cortex

- ▶ Cajal Ramon (the father & the mother of modern neuroscience)
- ▶ Pyramidal cells (1888)

